

Teleprompter & Video Creation App

A scripted-video recording tool for creators and teams.

Android first. Script generation, teleprompter capture, auto-captions, and smart clipping in a single release.

PLATFORM

Android
iOS in a later release

SCOPE

12 feature areas
Single V1 release

ESTIMATE

22–27 weeks
Solo build

AUDIENCE

Creators & teams
Realtors, coaches, corporate comms, course creators, sales

MONETISATION

Freemium + seats
Watermark loop · team billing

CORE STACK

CameraX · Media3
Supabase · cloud ASR

Document Structure

01	Product Decision	<u>3</u>
	Scope, wedge, and the two moat mechanisms	
02	Feature List	<u>4</u>
	Twelve areas, all in V1	
03	Technical Stack	<u>10</u>
	Choices, exclusions, and unit economics	
04	Excluded From Scope	<u>11</u>
	What was cut and why	
05	Build Order	<u>12</u>
	Dependency chain	
06	Timeline	<u>12</u>
	Estimates and the ranked cut list	
07	Standing Principle	<u>13</u>
	The test every scope decision must pass	

Product Decision

What the product is, what it deliberately is not, and the structural reason it can defend a position in a commodity category.

1.1 Scope

Teleprompter capture, auto-captions, and smart clipping, with AI script generation as the entry point. Android first; iOS follows once the composition pipeline is proven.

1.2 The rejected alternative

The slideshow-maker concept was dropped. The category is saturated — CapCut, VN, and Google Photos Memories all occupy it, the last one free and pre-installed. Competition there is won by paid user acquisition rather than by product quality, and consumer users churn off subscriptions within weeks. A slideshow mode can reuse the same Media3 composition layer as a second entry point in a later release if it ever becomes worth building.

1.3 The wedge

No general-purpose editor owns the *recording* moment for scripted talking-head video. Every competitor works on footage that already exists. A script read aloud while the camera runs is a workflow CapCut does not replicate, and it produces structured data — the script, the transcript, and word-level timing — that nothing else in the category has access to.

1.4 Moat mechanisms

Mechanism 01

Watermark distribution loop

Free-tier exports carry a subtle watermark. The audience is creators publishing to public feeds, so every single export becomes a distribution event rather than a private file. Acquisition cost trends toward zero as the free base grows.

Mechanism 02

Team scripts

A marketing lead writes and assigns scripts; sales reps record them. The billable unit becomes the account rather than the individual, which changes the churn profile entirely — a rep leaving does not cancel the subscription.

THIRD MECHANISM — ACCRUES OVER TIME

Brand-kit context fed into AI script generation means a team's outputs get progressively more on-brand with use. Approved scripts become tone samples for future generations. This is a switching cost built out of accumulated data, and it is the one thing a commodity teleprompter app cannot copy by shipping a feature.

Feature List

Twelve areas, all scoped into the V1 release.

2.1 Script Generation (AI)

Solves the blank-page activation problem. Every other feature in this document assumes a script already exists — without generation, a first-run user opens an empty text box and the teleprompter is useless until they write 200 words.

- Input: topic or rough notes
- Controls: target duration (15s / 30s / 60s / 2min), tone (direct, warm, energetic, formal), audience
- Output: 2–3 variants per generation, never one — picking beats accepting, and it hides the times the model misses
- Word count targeted to duration at ~140 wpm, reconciled against actual teleprompter scroll speed so the estimate is not a lie
- Regenerate a single paragraph without regenerating the whole script
- Inline rewrite on selected text: shorter / punchier / more formal
- Hook-first structure enforced in the prompt — the first sentence must earn the next three seconds
- Org brand-kit injected as context: company description, product, audience, banned phrases, tone samples from previously approved scripts
- Metered on the free tier

PROMPT CONSTRAINTS – NON-NEGOTIABLE

Short sentences. Second person. Contractions required. No corporate filler — no *"in today's fast-paced world,"* no *"unlock,"* no *"elevate."* Generic LLM copy read aloud on camera sounds exactly like generic LLM copy, and that damages trust in the entire app. Test the output against a real 30-second read before shipping. If it does not sound like a person talking, the feature is a liability rather than a differentiator.

2.2 Script Import

- Paste plain text
- URL fetch from a blog post or landing page, with a **"turn this into a 60-second script"** conversion step — not raw verbatim insertion, because nobody reads an article aloud
- File import via Storage Access Framework: `.txt`, `.docx`
- Google Docs link
- Repurpose an existing script to a different length or platform

2.3 Teleprompter & Recording

- Script editor: type, paste, or import
- Scrolling overlay on the live camera preview
 - Adjustable scroll speed, font size, text opacity, background scrim opacity
 - Tap-to-pause and resume; manual drag to reposition within the script
- Mirror mode for external teleprompter glass
- Countdown timer before recording starts
- Front / rear camera switch mid-session
- Resolution: 1080p / 720p
- Aspect ratio: 9:16, 1:1, 16:9, with a guide frame shown in preview
- Manual exposure lock and slider
- Tap-to-focus
- Live audio level meter and headphone monitoring
- Multiple takes per script, stored together, with a take picker and preview scrub
- Auto-save on interruption — incoming call, app backgrounded
- Keep-screen-awake during recording

2.4 Beauty & Camera Effects **OFF BY DEFAULT**

User-adjusted only, never automatic. This removes the *"why does my face look wrong"* first-run problem, keeps the preview pipeline at zero shader cost until a slider is touched, and preserves trust with corporate users who actively distrust visible facial modification.

- Master toggle, off on first launch, state remembered per user
- Skin smoothing: 0–100 slider, face-region masked via ML Kit face contours so the background stays sharp
- Warmth / tone: 0–100 slider
- Brightness lift ("ring light"): 0–100 slider
- Colour grade LUTs: 6–8 presets, with **None selected by default**
- Live preview of every adjustment before recording
- Reset-to-default button
- Effects applied at export rather than baked into the recorded file — changeable after the take is shot, so a bad filter choice never costs a take

EXCLUDED

Face reshaping, jaw slimming, and eye enlargement. Mesh warping against a face mesh is the expensive part of the pipeline, and it is wrong for this audience — a realtor wants to look rested, not altered.

2.5 Captions

- Auto-generate with word-level timestamps
- Full transcript editor — fix names, jargon, and ASR errors
- Word-level timing adjustment for drift correction
- Five themes covering position, font, size, and highlight style (karaoke word-highlight, pop-on, single-line)
- Brand colour picker applied to the caption highlight
- Logo upload with position and size control
- Per-caption manual repositioning
- Toggle captions off entirely
- English at launch

2.6 Editing

- Trim head and tail on a waveform-backed timeline
- Split and delete interior sections
- Reorder takes into a single sequence — multi-take stitching
- Frame-accurate scrub preview
- Undo / redo

2.7 Smart Clipping

The strongest differentiator in the release. Competitors run ASR against a video they know nothing about and guess at a hook. This app has the script the user read, word-level timestamps, and the intended structure — ground truth nobody else in the category has.

- **Filler-word removal.** Detects "um," "uh," "like," and "you know" from word timestamps. Each appears as a reviewable strike-through in the transcript, with remove-all or remove-individually.
- **Silence trimming.** Adjustable threshold, default 700ms.
- **Segment extraction.** An LLM analyses the transcript and returns ranked, self-contained clips with hooks, cut on word boundaries. The user picks which to export.
- **Target clip length.** 15s / 30s / 60s.
- **Auto-reframe.** Landscape to 9:16 with face tracking and a low-pass-filtered pan to prevent jitter.
- Manual override of the reframe crop position
- Every automated cut is shown in the transcript editor and is reversible before export

2.8 Audio

- Import own tracks via Storage Access Framework (`audio/*`)

- Independent volume mixing for the voice track and the music bed
- Auto-ducking of music under speech at -12 to -15 dB, located using caption timestamps rather than signal processing
- Fade in and fade out on the bed
- Trim the music track to video length
- Mute original audio option
- Graceful, plain-language failure for DRM-protected files from Spotify and Apple Music — otherwise it reads as a crash

NO BUNDLED MUSIC LIBRARY

"Commercial-free music beds" means a licensed catalogue. Shipping unlicensed tracks transfers Content ID strike liability to the user, and then back to the app.

2.9 Export & Share

- Export to gallery via MediaStore
- Quality selector: 1080p / 720p
- Progress indicator with a time estimate, backgroundable, with a completion notification
- System share sheet — `ACTION_SEND` with FileProvider
- Export the same video at multiple aspect ratios in a single pass
- Export the transcript as `.srt` or `.txt`
- Export history — re-export a project without redoing the edit

EXPECTATION SETTING

Effects force a full re-encode. A ten-minute 1080p export on a mid-range Snapdragon takes minutes, not seconds. Real progress reporting and backgrounding are mandatory, or the reviews will say "*freezes.*"

2.10 Free Tier & Watermark

FREE • UNLIMITED	Teleprompter recording, local takes, trim, beauty effects, audio import
FREE • METERED	Caption minutes (starting at 10 per month) and AI script generations (a handful per month)
FREE • WATERMARK	Subtle corner mark on every export — small, semi-transparent, positioned to survive platform crops on 9:16
PAID	Removes the watermark, raises caption minutes and the generation cap
PRICING	Monthly and annual only. No weekly plan — Play scrutinises weekly-plus-trial patterns, and it is the single biggest driver of one-star reviews in this category.

Two independently metered resources give two soft paywall triggers alongside the watermark, so conversion does not depend on a single moment of friction.

2.11 Accounts & Teams

- Email and Google sign-in
- An org is created on signup — a solo user is simply an org of one
- Roles: owner, writer, recorder
- Shared script library with folders
- Assign a script to a specific recorder, with a push notification on assignment
- The recorder's finished video is visible to the writer, with approve / request-retake
- Comments on a submitted video
- Org-level brand kit — logo, colours, default caption theme — inherited by every member and fed to AI script generation
- Seat-based billing with invite-by-email-link
- Script version history

2.12 Cross-Cutting

- Cloud sync of scripts and project metadata; video files stay local until explicitly uploaded for review
- Offline recording, with caption generation queued until connectivity returns
- Crash and ANR reporting
- Export analytics — which features are used before a paid conversion

DAY-ONE ARCHITECTURAL REQUIREMENT

Org-scoped data model with `org_id` on every row, from the very first migration. Retrofitting multi-tenancy after 500 solo accounts have signed up is a migration to avoid rather than survive. Supabase with RLS policies on `org_id` gets there fast.

Technical Stack

Selections, deliberate exclusions, and the cost constraint that pricing is designed around.

LAYER	CHOICE	RATIONALE
Camera	CameraX	<code>CameraEffect</code> (1.3+) allows OpenGL shader injection into both preview and capture
Composition & export	Media3 Transformer	<code>Composition</code> plus <code>EditedMediaItemSequence</code> for trimming; <code>OverlayEffect</code> with <code>BitmapOverlay</code> for captions and logo; <code>AudioProcessor</code> for mixing. Hardware-accelerated, first-party, actively maintained
Caption rendering	Canvas → Bitmap → time-ranged overlay	Full styling control with a straightforward implementation path
Beauty shaders	GLSL, written once	The same source is used as a <code>CameraEffect</code> for live preview and a Media3 <code>GLEffect</code> at export. Record clean, bake at export — fully non-destructive
Face detection	ML Kit face contours	Free and on-device; also drives auto-reframe subject tracking
ASR	Cloud — Deepgram or AssemblyAI	Both return word-level timestamps natively, which highlight captions and smart clipping both require. Roughly \$0.25–0.40 per audio hour
LLM	One integration, two prompts	Script generation and segment extraction share the same client
Backend	Supabase with RLS on <code>org_id</code>	Fastest path to correct multi-tenancy

3.1 Explicitly not used

TECHNOLOGY	REASON
FFmpegKit	Retired in early 2025 with prebuilt binaries pulled from Maven. Building on it means self-compiling and inheriting the LGPL/GPL question on a commercial app
On-device Whisper	Too slow on the mid-range Android hardware this audience actually owns
Third-party beauty SDKs	DeepAR and Banuba run several thousand per year — not justified at this stage when the shader work is a few weeks

UNIT ECONOMICS

ASR cost is the constraint to price around: **captions are metered, the teleprompter is not**. Script generation is cheap per call but will be hammered, so it is metered separately rather than folded into the same allowance.

SECTION 04

Excluded From Scope

Features present in the reference apps that will not be built, with the reasoning recorded so the decisions do not get relitigated.

FEATURE	REASON FOR EXCLUSION
AI eye-contact correction	Research-grade real-time gaze redirection. NVIDIA ships it on desktop GPUs. Not viable on a Snapdragon 6-series
One-click Instagram posting	No public content-publishing API exists for personal accounts. Every competitor's version of this is the share sheet with different marketing copy
One-click YouTube posting	The Data API upload works, but requires an OAuth verification review and costs 1,600 quota units per upload against a 10,000 daily default. Deferred
Bundled music library	Licensing liability. Users import their own audio instead
Chroma key (green screen)	ML Kit Selfie Segmentation covers the same job for a wider audience without requiring a green screen. Deferred, not dead
Cloud sync & web composer	BIGVU's enterprise layer, built over years. Not a V1 concern
Weekly subscription	Play Store scrutiny plus predictable review damage
Face reshaping	Wrong for the audience and expensive to build
Slideshow maker	Saturated category won on ad spend rather than product. Possible later reuse of the same composition layer

Build Order

The dependency chain is not negotiable even though the release is a single V1. Each step below depends on the one before it.

- 1

Camera, teleprompter, and takes
CameraX preview with the scrolling overlay, take storage and picker
- 2

Export pipeline
Media3 Transformer — the GL pipeline that steps 5 and 6 both depend on
- 3

Captions
ASR integration producing word-level timestamps
- 4

Transcript editor
Correction UI; the surface every automated cut is reviewed in
- 5

Smart clipping
Depends entirely on steps 3 and 4 being solid. Filler removal, then segments, then auto-reframe
- 6

Beauty shaders
Depends on the GL pipeline established in step 2
- 7

Auth and orgs
Org-scoped schema was already in place from the first migration; this turns on the UI
- 8

AI script generation
A generic version can ship earlier for testers, but the brand-kit context from step 7 is what makes it good
- 9

Team assignment and review loop
Assignment, notifications, approve / request-retake
- 10

Billing
Seat-based, monthly and annual

Timeline

Solo-build estimates by component, and the order in which scope should be cut if the schedule runs long.

COMPONENT	ESTIMATE
Teleprompter and captions core	8-10 weeks
Smart clipping	included above

COMPONENT	ESTIMATE
Beauty shaders	2–3 weeks
Auth, orgs, script sync	3–4 weeks
AI script generation and import	2–3 weeks
Audio import	< 1 week
Total, solo build	22–27 weeks

6.1 If scope must be cut

In this order. These cut cleanest without breaking anything downstream.

- 1 Multi-take stitching**
 The take picker alone is sufficient for the core workflow
- 2 Approve / request-retake review loop**
 Script assignment alone still makes the org the billing unit, which is the point
- 3 Beauty filters**
 Table stakes rather than a reason anyone switches apps

SECTION 07

Standing Principle

THE TEST

The build is not the hard part. Teleprompter apps are a commodity and there are dozens already on Play. **The moat is the watermark distribution loop, the team account structure, and the accumulating brand-kit context** — not the feature list.

Every scope decision in this document, and every one made after it, should be tested against whether it strengthens one of those three. A feature that strengthens none of them is a feature that makes the app better and the business no more defensible.